

The Subjective Moral Duty to Inform Oneself before Acting

Author(s): Holly M. Smith

Source: *Ethics*, Vol. 125, No. 1 (October 2014), pp. 11-38

Published by: The University of Chicago Press

Stable URL: <http://www.jstor.org/stable/10.1086/677024>

Accessed: 03-04-2018 22:46 UTC

JSTOR is a not-for-profit service that helps scholars, researchers, and students discover, use, and build upon a wide range of content in a trusted digital archive. We use information technology and tools to increase productivity and facilitate new forms of scholarship. For more information about JSTOR, please contact support@jstor.org.

Your use of the JSTOR archive indicates your acceptance of the Terms & Conditions of Use, available at <http://about.jstor.org/terms>



JSTOR

The University of Chicago Press is collaborating with JSTOR to digitize, preserve and extend access to *Ethics*

The Subjective Moral Duty to Inform Oneself before Acting*

Holly M. Smith

The requirement that moral theories be usable for making decisions runs afoul of the fact that decision makers often lack sufficient information about their options to derive any accurate prescriptions from the standard theories. Many theorists attempt to solve this problem by adopting subjective moral theories—ones that ground obligations on the agent's beliefs about the features of her options, rather than on the options' actual features. I argue that subjective deontological theories (such as Prichard's and Ross's) suffer a fatal flaw, since they cannot appropriately require agents to gather information before acting.

As decision makers we humans are epistemically far from perfect: we all too frequently lack accurate and complete information about the options we must choose among. Hence on many occasions we are unable to use the standard moral theories to guide our decisions, since we lack the information required to do so. This is notoriously true of act utilitarianism but is true for many traditional deontological theories as well.

Moral philosophers have advocated various responses to this situation, but one prominent response is to argue that it is a necessary condition for a moral theory that it be usable by actual decision makers to guide their decisions, however faulty or impoverished the decision maker's in-

* Work on this article was supported by a visiting fellowship at the Australian National University School of Philosophy and by the Carl and Lily Pforzheimer Foundation Fellowship at the National Humanities Center. I am grateful for this support, as well as for comments from the hardworking referees on earlier versions; from audiences at the University of Toronto, Fordham University, Notre Dame University, the City University of New York, Rutgers University, the University of North Carolina at Chapel Hill; and insightful comments by Nick Beckstead.

Ethics 125 (October 2014): 11–38

© 2014 by The University of Chicago. All rights reserved. 0014-1704/2014/12501-0001\$10.00

formation. Since the traditional moral theories fail to meet this condition, it is often urged that in their place we should adopt subjectivized versions of these theories. These versions make the fundamental moral status of an action depend, not on the objective features of the agent's circumstances or the objective consequences of the action, but rather on the beliefs (or credences) the agent has about the relevant features of the act, beliefs that are more readily accessible than facts about the world external to the agent. Thus many contemporary utilitarians advocate a version of utilitarianism according to which an agent is not obliged to maximize actual utility, but rather to maximize expected utility, where the expectation of a certain outcome depends on the agent's subjective estimate of its probability.¹

Deontologists have paid less attention to the epistemic difficulty that decision makers may have in applying deontological codes, perhaps relying on their sense that deontic duties are less subject to epistemic problems than consequentialist obligations. Indeed, it is typically easier to know when one is keeping a promise than when one is maximizing welfare. Nonetheless epistemic difficulties arise in applying deontological codes, just as they do for consequentialist codes. One may forget what one promised to buy at the supermarket; one may be unsure how much damage one has caused, and thus unsure what compensation is appropriate; or (as a school principal) one may be unsure whether an accused student was the one who bullied a younger student and thus unsure whether he deserves to be penalized. Hence deontologists, like consequentialists, must determine how deontological codes should accommodate false or incomplete beliefs on the part of decision makers. The subjectivizing strategy may prove attractive for them, as it has for consequentialists. Indeed, the most prominent early advocates of the subjectivizing strategy for accommodating ignorance and mistake are two deontologists: H. A. Prichard and W. D. Ross. These theorists famously rejected standard objective moral codes in favor of subjective codes which ascribe rightness and wrongness in light of the agent's beliefs about her situation.²

There are many issues to be raised about this "subjective turn" in normative theory. Elsewhere I argue that subjective moral theories do not in fact provide a fully effective solution to the problem of imperfect information.³ But in this article I will focus on another issue for such codes. Let us call a "subjective moral code" a code that makes an agent's primary

1. A leading example is Frank Jackson in "Consequentialism and the Nearest and Dearest Objection," *Ethics* 101 (1991): 461–82, 464–65.

2. H. A. Prichard, "Duty and Ignorance of Fact," in *Moral Obligation and Duty and Interest* (Oxford: Oxford University Press, 1968), 18–39; and W. D. Ross, *Foundations of Ethics* (Oxford: Clarendon, 1939), chap. 7.

3. Holly M. Smith, "Subjective Rightness," *Social Philosophy and Policy* 27 (2010): 64–110.

duty⁴ depend on her beliefs or credences about the nonnormative features of her action rather than on its actual nonnormative features.⁵ The question on which I shall focus is whether such a code can underwrite an appropriate obligation to gather information about one's future acts before deciding which act to perform. I take it as a datum that we all believe that there often is such an obligation, and a very strong one.⁶ If you are a military leader whose lieutenants recommend bombing a compound that might house enemy soldiers, you have an obligation to investigate—before bombing it—whether the compound really does house enemy soldiers, and whether it houses innocent civilians as well. If you are backing out of your garage, you have a duty to check your rearview mirror before proceeding. If you are a faculty member considering how to vote in an upcoming tenure case, you have an obligation—before voting—to inform yourself about the achievements of the tenure candidate.

Some philosophers have raised questions about whether a subjective moral code can underwrite such obligations to inform oneself before acting. Prichard himself took this to be a significant (although answerable) worry about the subjective code he advocated, and subsequent authors have raised the same question.⁷ In this essay I will examine this question as it arises for deontological codes. In the case of objective deontological codes, I will examine two different strategies for accommodating the obligation to gather information. The first strategy attempts to define a special, freestanding deontic duty of this type. I will argue that this strategy is not promising and advocate an alternative approach, namely, defining a derivative obligation to gather information in order to ensure that one satisfies one's subsequent deontic duties. Although this approach appears successful in the case of objective deontological codes, I will argue that neither it nor the alternative "freestanding duty" approach is successful in the case of subjective deontological codes. Thus Prichard was right to worry: it appears that subjective deontological codes cannot underwrite an obligation to gather information before acting. Even if one holds that subjective deontological codes provide a solution to the problem of imper-

4. I use the term "primary duty" because there are "hybrid" moral theories that recognize a primary objective account of duty, but also a "secondary" subjective duty to perform the best act relative to the agent's beliefs.

5. Note that this definition does not make her subjective duty depend on her belief that she has a duty.

6. In this article I will use "one has a duty," "one ought," and "one has an obligation" to mean the same thing.

7. Prichard, "Duty and Ignorance of Fact," 27–28. See also Bart Gruzalski, "Foreseeable Consequence Utilitarianism," *Australasian Journal of Philosophy* 59 (1981): 163–76, 169; Roy Sorensen, "Unknowable Obligations," *Utilitas* 7 (1996): 247–71, 254; and Michael Zimmerman, *Living with Uncertainty* (Cambridge: Cambridge University Press, 2008), 14.

fect information, one should reject such codes on the ground that they cannot include a suitable obligation to investigate before acting.

I. FOUR CAVEATS

Several caveats before we begin. First, I will focus primarily on agents who have false beliefs or no relevant beliefs, rather than agents who are merely uncertain about their prospective action's character. This restriction will not distort the inquiry or its outcome. If a moral theory flunks the test of appropriately handling the duty to gather information for agents having false beliefs or no relevant beliefs, then it is inadequate, even if it passes this test for agents laboring under uncertainty.

Second, we will inspect subjective moral codes that prescribe actions in light of the agent's beliefs—her actual beliefs. But some may think that a person's duty depends on what it would be reasonable for the person to believe, not on what the person actually believes.⁸ These theorists would argue that we should concentrate on reasonable-belief versions of these codes, which may elude the problems I raise for actual-belief versions.

In other work I argue that reasonable-belief versions of subjective codes are even less effective than actual-belief codes at serving the purpose for which they have been introduced, namely, to provide moral codes that are usable by any agent, however deficient his information may be.⁹ Given this, my primary focus will be on actual-belief codes. Nonetheless I will show later that prominent reasonable-belief versions of subjective codes are subject to the same problems that I shall describe for parallel actual-belief versions.

Third, although I will examine situations in which agents have false beliefs about nonnormative facts (such as what one promised to buy at the supermarket), I will assume that these agents all have correct beliefs about general normative facts (such as whether or not violating a promise is *prima facie* wrong). Of course agents make mistakes about normative facts as well as about nonnormative facts, but dealing with that complication must be left for another occasion.

Fourth, it might be hoped that epistemologists have investigated the nature of our purely epistemic duty to acquire information and can provide guidance in formulating a parallel moral duty. Unfortunately little

8. Advocates of the "reasonable belief" view are legion. Examples include Brad Hooker, *Ideal Code, Real World* (Oxford: Oxford University Press, 2000), sec. 3.1; Zimmerman, *Living with Uncertainty*; Shelly Kagan, *The Limits of Morality* (Oxford: Clarendon, 1989), 1; and Thomas Scanlon, *Moral Dimensions: Permissibility, Meaning, Blame* (Cambridge, MA: Belknap Press of Harvard University Press, 2008), 47ff.

9. Holly Smith, "Subjective Rightness," and "The Moral Clout of Reasonable Beliefs," in *Oxford Studies in Normative Ethics*, vol. 1, ed. Mark Timmons (Oxford: Oxford University Press, 2011), 1–25.

substantive work has been done in this field, and many of the suggestions that have been made appear to be vulnerable to the same problem that I shall point out for subjective deontological moral theories. Thus I shall not be able to invoke helpful insights from epistemology.¹⁰

II. THE DUTY TO ACQUIRE INFORMATION BEFORE ACTION: OBJECTIVE DEONTOLOGICAL THEORIES

Our normal view is that when we must make an important choice in the future, we have a duty now to obtain the information that will reveal which option we ought to choose on that future occasion. Our question is whether subjective deontological codes can accommodate such a duty. I shall argue that they cannot. To see this, let's start by considering how such a duty arises within objective codes, codes in which the agent's obligations depend on objective facts about the world, not on what the agent believes about the world.

Consider an agent who must lay off an employee during a financial re-trenchment. Claire is a human resources manager, who tomorrow (Tuesday) must lay off one of the employees from her financially stressed company, and who will have a duty to carry out this task in the best manner possible. As of today (Monday) Claire knows relatively little about the two employees, Max and Mina, who are candidates for being laid off. However, she could spend time today gathering accurate information about the employees' length of service, productivity, number of dependents, and so on and if she did so, she would have enough information to make a good decision tomorrow. Normally we would think that Claire has a duty to gather this information before making her decision.

But why does Claire have this duty, within an objective moral code, to acquire information? If the code is a consequentialist one, requiring Claire to maximize human welfare, the answer is straightforward. Suppose Max has fewer dependents than Mina. It would maximize welfare—that is, mitigate negative impact on welfare—for Claire to lay off Max tomorrow. However, suppose that as things are, Claire believes that it would maximize welfare for her to lay off Mina, since she (falsely) believes Mina has fewer dependents. Nonetheless if she investigated she would come to believe (correctly) that it would maximize welfare to lay off Max rather than Mina. Thus, assuming that Claire is disposed to follow this code, it is true that if she investigated, she would subsequently lay off Max rather than Mina. Hence Claire's gathering information in this situation just is her maximizing welfare, in virtue of this act's leading to her subsequent act

10. One useful resource is Richard Feldman, "Epistemological Duties," in *The Oxford Handbook of Epistemology*, ed. Paul K. Moser (Oxford: Oxford University Press, 2002), 362–84. I will take up the question of purely epistemological duties in a separate paper.

of releasing the employee who would be least badly affected by losing the job. This example shows that objective welfare-maximizing codes can readily explain the duty to acquire information about one's future acts, because acquiring information is often just a way of bringing about the best consequences. Of course this is not always true: sometimes investigations can mislead the investigator or are so costly that the good consequences they would secure do not counterbalance the bad effects they occasion. But even in these cases an objective welfare-maximizing code provides a code-appropriate recommendation, precisely because it takes these factors into account in determining whether or not the agent has a duty to investigate before acting. We do not think, for example, that Claire ought to spend time investigating if doing so would distract her from responding to a law suit that would send her company into bankruptcy.

However, the situation is less clear-cut when the objective code incorporates non-welfare-maximizing elements, such as duties to keep promises, to compensate others for past wrongs, and to act justly. Let us call these "deontic duties" (and codes containing them "deontological codes"). Clearly we think that one can have a duty to gather information now in light of one's subsequent deontic duties. Just as one has a duty to investigate in order to ensure that one subsequently acts to maximize welfare, so, if one has backed one's truck into another person's car, one has a duty to investigate whether one has caused any damage, so that one can fulfill any resultant duty to compensate the car's owner. But although we can understand gathering information as a way of maximizing welfare, and hence as merely a special case of one's duty to maximize welfare, the same approach is not possible in the context of gathering information in order to ensure that one fulfills the deontic duties I described. Gathering such information is not a way of keeping a promise, or compensating others, or of acting justly.¹¹ If there is a duty to gather information in connection with these duties, a different account of it must be given.

11. Indeed, in calling these duties "deontic duties," all I mean to indicate is that these duties are ones about which it is not plausible to say that gathering information that leads to one's fulfilling such a duty itself counts as a special case of fulfilling the "deontic" duty. There are, of course, other duties traditionally recognized as "deontic," such as the duty not to injure others. However, these duties could arguably be thought of as essentially consequentialist in character—the duty not to injure others is a duty not to produce a certain kind of consequence, namely, harm to another person (or harm that the person could reasonably reject). Hence the duty to acquire information in order to ensure that one fulfills the duty not to injure others could potentially be justified as a way of not harming another person. Admittedly, this analysis is something of a stretch. However, to ensure that we are discussing the duty to acquire information in connection with fulfilling pure deontic duties that can't be conceptualized as types of consequentialist duties, I restrict my attention to the examples in the text. Almost all deontological theories recognize such duties, so some way of dealing with the acquisition of information relevant to them must be found.

Suppose Claire's moral code doesn't require her to maximize welfare in selecting which employee to lay off but instead requires her to act justly in making this decision. In particular it stipulates that justice requires her to first lay off those employees who are least productive. Mina is the least productive employee, so Claire will act justly only if she releases Mina. However, Claire falsely believes Max to be the least productive employee, so if she doesn't gather any additional information, she will lay off Max. On the other hand, if she were to gather more information, she would come to believe that Mina is the least productive employee and act justly by laying her off instead of Max.

Under this code, should Claire gather information before deciding whom to lay off? Clearly an adequate deontological code must incorporate a special duty requiring information to be gathered before taking action that may be just or unjust. What would be the content of such a duty? Some theorists, grappling with this issue, have tried to formulate such a duty as a freestanding deontic duty to gather information: freestanding in the sense that it is not defined or justified in terms of the consequences of gathering information for the agent's subsequent choice. Thus (if we convert Prichard's remarks about a subjective code into remarks about an objective code) Prichard suggests that

M1. One has a prima facie duty to consider the circumstances as fully as one can.¹²

If we convert Ross's remarks on this question to a suggestion about an objective code, we may interpret him as advising that

M2. One has a prima facie duty to investigate the circumstances relevant to an upcoming duty until either (*a*) one must act because time has run out or (*b*) one has reached the point where no further consideration would enable one to judge better of the circumstances relevant to the duty.¹³

More recently Gideon Rosen has asserted that "we are under an array of standing obligations to inform ourselves about matters relevant to the moral permissibility of our conduct: to look around, to reflect, to seek advice, and so on." Noting that "the content of this obligation varies massively from case to case," he gives as an example "I am under an obligation to look out for other people when I'm out walking."¹⁴ We could construe this advice as follows:

12. Prichard, "Duty and Ignorance of Fact," 27.

13. Ross, *Foundations of Ethics*, 157.

14. Gideon Rosen, "Culpability and Ignorance," *Proceedings of the Aristotelian Society* 108 (2002): 61–84, 63.

M3. One has a great variety of *prima facie* obligations to inform oneself about matters relevant to the moral permissibility of one's future conduct, such as an obligation to look out for other people when one is out walking.

Let us review these proposals, construing them as descriptions of duties that can be defined or justified without reference to their impact on the agent's subsequent actions. One virtue of these proposed duties is that an agent who lacks significant information about his upcoming act would nonetheless be able to follow a freestanding rule to decide whether or not to gather more information. But a review of the proposed rules suggests that they are uniformly ill conceived. Considering the circumstances as fully as one can, as Prichard suggests, may lead one to spend unnecessary time and effort on this endeavor, since one may already have enough information to incline one to perform what is one's subsequent duty, or one may obtain such information with very little effort, or further investigations won't improve one's information, or will even degrade it. Investigating until time has run out, as Ross suggests, has the same defects. Ross's suggestion that one should gather information to the point where no further consideration would enable one to judge better of the circumstances relevant to the duty is more promising. But it too may be misguided. There is an ambiguity in what Ross means by "judge better of the circumstances." If he means "gather information until one has achieved a true opinion about the matter," then this may be an impossible task, because there may be no available information that would lead to a true opinion.¹⁵ If he means "gather information until no further consideration would render one better justified in one's opinion," then it is not clear why one should expend the effort to achieve maximum justification, when less effort might lead one to the same belief about the matter, or when less effort might render one less well-justified but with a more accurate opinion (since one can be well justified in believing a falsehood). Nothing in Ross's description of this duty provides any guidance about what to do in cases where one should not aim at maximum justification. And Rosen's suggestion that a proper deontic code would include a raft of substantive rules to carry out certain kinds of investigations in different types of circumstances doesn't provide us what we need unless the content of those rules can be specified, a difficult challenge that Rosen doesn't undertake. Thus having a duty to "look out for people when one is out walking" may be reasonable when one is walking on a crowded street but not reasonable when walking on

15. I am assuming that Ross's duty to gather information would be a *prima facie* duty. It is difficult to see how the strength of such a duty could be settled without importing considerations about the impact of following the duty on one's future actions, making it no longer a freestanding deontic duty.

a mountain trail little trafficked by other hikers.¹⁶ Some theorists have suggested that there are many specific evidence-gathering duties that are attached to roles or jobs.¹⁷ Thus a surgeon has a duty to double-check which arm is diseased before performing an amputation, and a cashier has a duty to check a customer's signature before accepting her credit card. Such investigatory duties are often spelled out in institutional rules adopted to ensure that people occupying certain positions will secure the information they (or others) need in order to carry out their other duties effectively. But not all decisions whether to gather evidence are made by people insofar as they occupy one of these role-based positions, and we need to define the more general deontic duty to acquire information that applies in these broader circumstances. Rosen seems to be attempting this task, but his suggestions fall short.

Scrutiny of these proposed freestanding duties to seek information reveals that they have several flaws: either they provide poor advice, or they provide insufficiently detailed advice, or they provide advice for which no rationale is forthcoming.¹⁸ On reflection it seems clear that what a deontologist should want is a duty to gather information designed so that the agent who fulfills this duty will then fulfill (or will not violate) her subsequent substantive obligations. This is the rationale that seems to be inchoately behind some of the foregoing proposals, but they are not sufficiently sensitive to it. The proposed duties fail to take into account factors likely to influence the inquiry's effect on the agent's performance of subsequent duties, such as how much information he already has, how helpful the new evidence would be, and especially how significant the duties are whose satisfaction may or may not depend on his seeking additional information.

The most straightforward way to frame the duty to acquire information in a manner that takes this aim into account is for the duty to require gathering information just in case doing so would lead the agent to satisfy deontic duties subsequently incumbent on him to keep his prom-

16. No doubt a theorist attempting to define the freestanding duty to acquire information could try to formulate a suite of such rules, each rule designed to be used in a different situation. But then the question would arise why *these* rules rather than others? It appears that answering this question would either involve moving to a type of rule consequentialism (these rules would tend to produce the best consequences), or a type of rule deontology in which the rules are adopted because they do the best job of leading adherents to carry out their subsequent duties. The former rationale abandons deontology for a form of consequentialism, while the latter imports a version of the derivative deontic rationale for gathering information that I advocate later in the article.

17. This was suggested by members of an audience at Fordham University and by a referee for this journal.

18. Characterizations by virtue theorists of this duty are typically stated in such vague terms that they provide insufficiently detailed advice. See, for example, Jason Baehr, *The Inquiring Mind* (Oxford: Oxford University Press, 2011), 82–89.

ise, compensate others, or act justly. Let us call this nonfreestanding duty a “derivative” duty to gather information, since the fact that it is a duty derives from the subsequent duties it would lead the agent to satisfy. This duty will of course be a *prima facie* duty, since it may conflict with, and sometimes be overridden by, other more stringent duties incumbent on the agent.¹⁹

If Claire’s code included such a derivative duty, Claire would have a *prima facie* duty to gather information about her employees on Monday so that she would act justly in choosing on Tuesday which employee to lay off. It appears that a code containing a derivative duty of the sort described can provide an acceptable account of duties to gather information, precisely because it suitably grounds the duty to investigate on the effects of the agent’s fulfilling this duty on her carrying out her subsequent deontic duties.

However, such a derivative duty needs to be more nuanced. Although we have focused on cases in which the agent has the option of gathering information at time t_1 to inform her decision at later time t_2 , real life is more complex than this. Sometimes, for example, the information gathered at t_1 will be useful to more than one decision to be made later on. Thus a prosecutor might gather information at t_1 in order to inform his decision at t_2 about whether to bring charges against an alleged lawbreaker, but also to inform his subsequent decision at t_3 about what penalty to recommend if the defendant is found guilty. Moreover, these subsequent duties may have different degrees of stringency that will need to be weighed against each other. To accommodate these facts, let us introduce the notion of the “deontic value” of different deontic duties. The deontic value of an act expresses the weight, or stringency, of the duty to perform (or not to perform) that act (or, in other terminology, the force of the moral reason to perform or not to perform that action). Thus the deontic value of saving a person’s life is greater than the deontic value of keeping a minor promise. In determining what one ought all things considered to do, one weighs the deontic value of the various (sometimes conflicting) duties involved. If one has to choose between saving a life ver-

19. This formulation assumes that the derivative duty should be framed in terms of actualism rather than possibilism (in this context actualism says that our current duty depends on what duties it would actually lead us to fulfill, while possibilism says that our current duty depends on what duties it would enable us to fulfill). Possibilism seems ill equipped to generate appropriate duties to gather information. Since Claire (and many other agents faced with a decision whether or not to acquire more information) already is able—whether or not she investigates—to do the right thing in the future, a possibilistic version of the duty to investigate would not generate any objective duty for her to gather evidence. For my earlier discussions of these views, see Holly S. Goldman, “Dated Rightness and Moral Imperfection,” *Philosophical Review* 85 (1976): 449–87, and “Doing the Best One Can,” in *Values and Morals*, ed. Alvin Goldman and Jaegwon Kim (Dordrecht: Reidel, 1978), 186–214.

sus keeping a minor promise, one ought all things considered to save the life, since the deontic value of this act is greater than the deontic value of keeping the promise. This concept would need further refinement to support certain kinds of constraints that are often built into deontological theories, such as the constraint that it would be all things considered wrong for me now to kill a person even if killing the person ensures that you (or I myself) kill fewer people in the future. Since issues raised by these kinds of constraints are not implicated by the questions of this article, I shall not attempt to refine the concept of deontic value in order to reflect them.²⁰

Using the notion of deontic value, we are in a position to accept the following:

OD. An agent has an objective derivative prima facie duty to acquire information if and only if doing so would lead the agent subsequently to produce the maximum possible amount of deontic value (typically through his carrying out the various deontic duties that would later be incumbent on him).²¹

20. The notion of this sort of “deontic value” has been recognized, implicitly or explicitly, by a number of authors and seems especially crucial for any deontological attempt to state what an agent acting under uncertainty ought to do. See Goldman, “Dated Rightness and Moral Imperfection”; Robert Nozick, *Anarchy, State, and Utopia* (New York: Basic Books, 1974), 29n; Frank Jackson and Michael Smith, “Absolutist Moral Theories and Uncertainty,” *Journal of Philosophy* 103 (2006): 267–83; Shelly Kagan, “The Additive Fallacy,” *Ethics* 99 (1988): 5–31; John Broome, *Weighing Goods* (Oxford: Blackwell, 1991), chap. 1; Ted Lockhart, *Moral Uncertainty and Its Consequences* (New York: Oxford University Press, 2000), 80–97; and Jacob Ross, “Rejecting Ethical Deflationism,” *Ethics* 116 (2006): 754–55. Douglas Portmore defines an act’s “deontic moral value” as “a measure of how much [objective] moral reason there is to perform it” in *Commonsense Consequentialism* (Oxford: Oxford University Press, 2011), 19 n. 36. Some theorists doubt that a concept of quantified deontic value can be devised that appropriately reflects special deontic features, such as agent-centered options, supererogatory acts, and the self-other asymmetry (terminology from Portmore, *Commonsense Consequentialism*, 8), but I believe the jury is still out on this question. One of the lessons of this article is that a deontologist who rejects the notion of “deontic value” will be very hard-pressed to define a suitable duty to acquire information before acting.

21. In a case in which an agent has two or more possible investigatory acts at t_1 , I am conceiving of the derivative duty so that the agent has only one prima facie derivative duty at t_1 , namely, the “cumulative” prima facie duty to do whichever of those investigatory acts would produce the maximum deontic value at the later time. OD bears a family resemblance to what Alexander A. Guerrero calls “modest moral epistemic contextualism” in “Don’t Know, Don’t Kill: Moral Ignorance, Culpability, and Caution,” *Philosophical Studies* 136 (2007): 59–97, 70.

A referee for this journal questions the truth of OD in light of the following case: Jim, a demolition engineer, is obligated to see whether anyone is in a building which he has been hired to demolish. In investigating the building, he sees his partner having an affair. This leads Jim to commit suicide, leaving his younger orphaned sister to fend for herself, violating both a promise he made to his parents to watch out for her and a familial duty he had to care for her. In fact, at the time Jim would have demolished the building, everyone would already

This derivative duty is still merely *prima facie*, since it might conflict with, and even be outweighed by, a stronger duty to carry out some other deontic obligation at the very time the agent would investigate. For example, the only way to gather information may involve torturing an innocent person and so violate the more important duty not to torture. According to OD, Claire has a derivative duty to seek information on Monday before she releases one of the employees on Tuesday, since her investigating would lead to her successfully carrying out her later duty to act justly when she makes the layoff decision.

Our tentative conclusion should be that objective moral codes including deontic duties to keep promises, compensate others for past wrongs, and act justly, can include normatively appropriate duties to gather information. The duty to gather information relevant to these duties is a specially defined derivative *prima facie* duty that requires one to gather information when doing so ensures that one will maximize deontic value by carrying out one's subsequent deontic duties.

III. THE DUTY TO ACQUIRE INFORMATION: SUBJECTIVE MORAL THEORIES

I have described how objective deontological theories can successfully define a duty to acquire information before making a subsequent morally significant choice. But our main focus of interest is subjective deontolog-

have vacated. The referee contends that it is plausible that (a) Jim has a *prima facie* duty to investigate the interior of the building, and (b) Jim had an all-things-considered obligation to investigate the building, and (c) investigating did not maximize deontic value. He concludes that duty OD I describe in the text isn't correct. I don't agree that the case shows OD to be mistaken. The case is a complex one, with a number of factors influencing how we assess Jim's action. First, our intuition that Jim has a duty to investigate may heavily reflect our sense that Jim would be blameworthy for not investigating, given that he couldn't possibly foresee the actual consequences of his investigating. It's often hard to separate one's intuitions about what is objectively obligatory from what is blameworthy. Second, the case raises the possibility that the derivative duty to investigate should exclude subsequent duty-fulfillments that arise through what are sometimes called "deviant causal chains." (It is only by a deviant causal chain, not the normal causal chain we expect to link action to consequences, that Jim's investigation leads to his subsequent violations of his promises.) I will not try to explore this interesting possibility here. Third, I think that Jim's all-things-considered duty to investigate arises from an independent source, namely, his role-based obligations as a professional demolitions contractor. The duty characterized in the text is only meant to be a derivative *prima facie* duty that arises from the way in which the agent's investigating would affect his fulfillment of his subsequent duties. Such duties can conflict with other kinds of duties and sometimes be outweighed by them. Thus in this case Jim's *prima facie* derivative duty not to investigate may be outweighed by his role-based *prima facie* duty to investigate. Finally, the case may provide evidence that the deontological duty to investigate should be time-centered.

ical moral codes, not objective ones. Given our conclusions about the duty to acquire information in objective codes, what can be said about subjective codes? I will examine two different ways in which a subjective deontic code might define the duty to gather information and argue that neither of them deliver the kind of duty to gather information that we believe is appropriate.

A. *The Duty to Gather Information in Subjective Codes That Include Deontic Duties: Freestanding Duties*

The first approach to defining such a subjective duty to gather information (suggested by the parallel approach to the objective duty) would be to characterize this duty as a freestanding duty, not defined or justified by reference to the effect that the agent believes satisfying this duty would have on her fulfilling her subsequent deontic duties. But reasoning similar to our previous deliberations about objective freestanding duties to gather information suggests that the strategy of invoking such a freestanding duty is unlikely to generate a normatively acceptable subjective duty. Just as any acceptable objective duty to gather information must take into account the effect of gathering information on the agent's subsequent actions, so any acceptable subjective duty to gather information must take into account the effect the agent believes the act would have on her later actions. The kinds of freestanding duties to gather information suggested by Prichard, Ross, and Rosen, when translated into subjective terms, fail this test. For example, a subjective duty to do what you believe is considering the circumstances as fully as you can is not a satisfactory account of this duty. Why should you do what you believe is considering the circumstances as fully as you can, if you believe that doing so will commit you to spending far more effort in gathering information than is necessary in order to arrive at the truth? Similar remarks can be made about the subjective versions of Ross's suggestions. And subjective versions of Rosen's suggestions lack the content and rationale to constitute a satisfactorily substantive proposal. Thus trying to define a subjective duty to gather information as a freestanding duty appears to be the same blind alley that it was in trying to define an objective duty to gather information.

B. *The Duty to Gather Information in Subjective Codes That Include Deontic Duties: Derivative Duties*

Since the "freestanding" duty approach to defining the subjective duty to gather information appears wrongheaded, let us turn to a second approach. Modeled after the second approach to defining the parallel objective duty, this approach defines the subjective duty to gather information as a derivative duty. As we saw before, this derivative duty must reflect

the fact that the agent may believe that her informing herself now may have an impact on her carrying out *multiple* deontic duties at later times, and that these duties may have different degrees of stringency. In light of this, we can propose the following:

SD. An agent has a subjective derivative prima facie duty to do what he believes is acquiring information if and only if he believes that doing what he believes is gathering information would lead him subsequently to produce the maximum amount of deontic value (typically through his doing what he then believes to be carrying out the various deontic duties that would then be incumbent on him).²²

Just as in the case of objective duty, this appears to be a natural way to define the derivative subjective duty to gather information. It, too, is merely a prima facie duty, since the agent may realize that the process of gathering information may violate some more important deontic duty (such as the duty not to engage in what he believes is torture).

To see how this approach would work, suppose Claire, the human resources manager tasked with laying off an employee, is governed by the following subjective Code C, which includes a prima facie deontic duty to do what the agent believes to be laying off employees justly and a derivative prima facie duty to do what the agent believes to be gathering information before acting. Justice will be best served by doing what the agent believes to be laying off the least productive employee, but in a case of employees who are all tied for minimum productivity, or a case in which the manager does not know which employee is least productive, the second-best solution called for is to lay off the person the agent believes to be the most recently hired employee.

Code C

- (1) When an agent believes she has been tasked to lay off employees in light of a financial retrenchment, the agent subjectively ought prima facie either to
 - (a) do what she believes to be laying off an employee if she believes that person to be the least productive employee, or

22. This represents our understanding of the duty. The agent himself need not conceive it in these terms. For example, Claire would fulfill this duty if she thinks to herself, "I should investigate because I estimate that doing so would result in my treating my employees most fairly when I choose whom to lay off." I also wish to leave open the possibility of stating SD more broadly so that it recognizes agents as fulfilling this duty if they conceive their future acts, not in terms of such explicitly moral terms as "duty" or "justice," but rather in terms of the underlying nonnormative features of the future acts that make them just. Some theorists might want to argue that SD (and OD) should be stated in a non-maximizing form, e.g., as satisficing principles.

- (b) do what she believes to be laying off the most recently hired employee if she does not believe of any employee that he or she is the least productive employee.
- (2) When an agent believes she has the opportunity to gather information relevant to a subsequent duty, then the agent subjectively ought *prima facie* to do what she believes to be gathering relevant information before fulfilling her subsequent duties if and only if she believes that doing so would subsequently lead her to produce the maximum amount of deontic value.²³

Code C contains an apparently plausible subjective derived *prima facie* duty to do what the agent believes to be gathering information relevant to her subsequent deontic duties. But does such a code impose actual appropriate duties to acquire information? Would agents governed by Code C ever be obligated—even *prima facie*—to gather information? To see the answer to this, we need to assign deontic values to Claire's carrying out, or violating, her various substantive deontic duties. Let's assume that the duty she would have to select the most recently hired employee for layoff if she doesn't believe of any employee that he or she is least productive constitutes a lesser form of justice than her duty to lay off an employee whom she believes to be least productive—it provides a second-best solution. Thus fulfilling this duty has less deontic value than fulfilling the duty to lay off the least productive employee. Plugging in plausible numbers for the deontic values of her acts, we get the representation of Claire's options shown in figure 1.²⁴

In this case, Claire currently doesn't have any beliefs about which employee is least productive, although she (correctly) believes that one of them is more productive than the other. She already believes that Max is the most recently hired employee. Claire truly believes that if she does what she believes is gathering information she will come to accurately believe of one of the employees that he or she is least productive, whereas if she does what she believes to be omitting to gather information, she will continue not to believe of any employee that he or she is least productive. If she gathers information, she will have a subsequent duty (by clause 1 *a* of Code C) to do what she believes to be laying off the least productive employee, whereas if she omits to gather

23. For brevity I have not stated the obvious clauses dealing with *prima facie* wrongness and permissibility. Again, the agent may use different concepts to conceptualize Code C.

24. To avoid overcounting, the acts of investigating (and not investigating) receive no deontic value of their own in fig. 1. Once it is established whether or not there is a *prima facie* derivative duty to investigate, acts of fulfilling or violating this duty would be assigned deontic values in order to be weighed against other potentially conflicting duties.

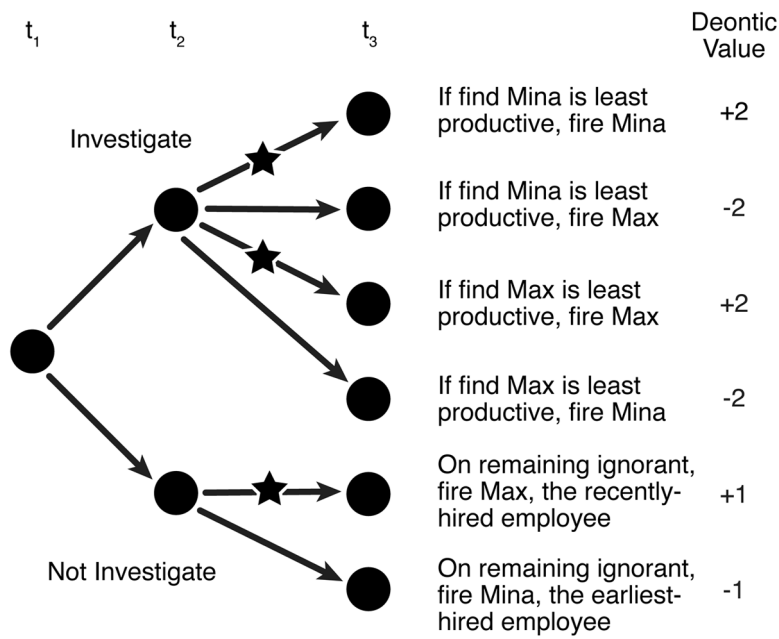


FIG. 1

information, she will have a subsequent duty (by clause 1 *b*) to lay off Max, whom she believes to be the most recently hired employee. (The stars in fig. 1 represent which choice Claire would actually make, depending on what information she has at t_2 after she either investigates or fails to investigate.) Claire truly believes that she will carry out whichever of these subsequent duties is incumbent upon her.

From figure 1 we can see that Code C tells Claire to do what she believes to be gathering information, since she believes that her investigating will create a duty either to lay Mina off or to lay Max off (depending on what she discovers about their productivity), either of which has a value of +2 when she carries it out. By contrast, she believes that doing what she believes to be not gathering information will create a weaker duty (to lay off Max, whom she believes to be the most recently hired employee) having a value of merely +1 when she carries it out. (Note that if she fails to investigate, she then has no duty to do what she believes to be laying off the least productive employee.) She can maximize deontic value by doing what she believes to be gathering information, and according to clause 2 of Code C this is what she should do. This is indeed the duty we want to ascribe to Claire in this situation; she should do what she believes to be gathering information. Thus this approach to the sub-

jective duty to gather information seems to ascribe the duty to investigate in a normatively appropriate way in Claire's case.

C. A Problem for This Approach, and the Solution

But further reflection reveals a problem for this approach. Claire's code is a subjective code. This means that her actions at t_2 of investigating (or not investigating) create her later duties at t_3 , since what duty she has at t_3 depends on what she believes at t_3 , and what she believes at t_3 in turn depends on whether or not she has investigated at t_2 . Her doing-what-she-believes-to-be-investigating at t_2 creates a duty at t_3 to do-what-she-believes-to-be-laying-off-the-least-productive-employee, while her doing-what-she-believes-to-be-not-investigating at t_2 creates a later duty at t_3 to do what-she-believes-to-be-laying-off-the-most-recently-hired-employee. Code C implies that she ought to do-what-she-believes-to-be-investigating at t_2 because that will create a more weighty duty (worth +2 rather than a mere +1) which she will then fulfill at t_3 . But is it really true that one should act so as to create a weightier duty for oneself which one would then carry out, as opposed to creating a less weighty duty for oneself which one would then carry out? This assumption lies behind our "derivative duty" approach to the duty to gather information, but it may be mistaken.

We can get some grip on this question by considering promises, which are one prominent way, even under an objective moral theory, that a person can create duties for herself. If there is no independent moral reason to make any promise, should one make a weightier promise as opposed to making a less weighty promise? Suppose Devon receives two e-mail messages: one from Kate, who asks Devon to take care of Kate's cat next week while she is on vacation, and one from Fred, who asks Devon to take care of his goldfish next week while he is on vacation. Both friends say that if Devon can't do it, they can find another equally competent caretaker. On the other hand, if he promises, then the pet will only be fed if Devon himself does it. Should Devon make a promise to take care of Kate's cat, rather than make a promise to take care of Fred's goldfish, because the duty he would create to take care of the cat would be a weightier duty (weightier because it would be worse if he fails to carry out the promise) than the duty he would create to take care of the goldfish? This doesn't seem right: given that someone else would ensure that each creature would be cared for if Devon doesn't make the promise, there is no more reason for him to promise to take care of the cat than to promise to take care of the fish just because it would be more important to carry out the first duty (if it came into existence) than to carry out the second duty (if it came into existence).

In earlier work I argued that it is counterintuitive to judge that someone like Devon ought to promise to take care of the cat rather than the

goldfish because it would give rise to a weightier duty, and that it is even more counterintuitive, in certain kinds of cases, to conclude, on the basis of the positive value that keeping a promise would add, that an agent has a duty to make a promise that would convert an otherwise wrongful act into an obligatory one. I then argued that the best way to avoid these counterintuitive implications is to hold that, while breaking a promise has negative value, keeping a promise as such has no positive value—no value above and beyond what the promised act would have had if it had not been promised.²⁵ On this view there is reason to avoid breaking a promise and also reason to avoid breaking a weightier promise rather than breaking a conflicting lesser promise. However, there is no obligation to make a weightier promise as opposed to a weaker promise, assuming one would keep whichever promise one made. Of course, there might be other, nondeontic reasons to make and then keep a weightier promise, but the fact that one promise is weightier than another does not in itself provide reason to make it. The disvalue of breaking a promise may imply that one ought not to make that promise in the first place. Thus suppose that if Devon promised to take care of the cat, he would then fail to do so, whereas if he promised to take care of the goldfish, he would carry out this duty. Since we ascribe negative value to Devon's violation of his duty to care for the cat (even if we don't ascribe positive value to his keeping his promise), in this case we can see that Devon ought to promise to take care of the goldfish, and then ought to do so, even though his possible duty to take care of the cat would be a weightier duty than his possible duty to take care of the goldfish.²⁶

The problems encountered by ascribing positive value to keeping a promise arise because making the promise creates a new duty. If we hold that one has a duty to create a new duty in order to garner positive value from subsequently fulfilling the newly created duty, we run into the kind of difficulties described in Devon's case. There are other kinds of cases as well in which a person's action may create a duty, which the person would then fulfill, but we don't conclude that the person ought to create the duty just so that he would act rightly in fulfilling it. Thus if a man impregnates a woman, he creates a duty to provide child care and support. Even if he would carry out this duty, we don't conclude from this that he has a *prima facie* duty to impregnate the woman and create the

25. Holly Smith, "A Paradox of Promising," *Philosophical Review* 106 (1997): 153–96.

26. In a critical article Earl Conee rejects this asymmetrical view about the values of keeping and breaking promises. But he does this at the price of denying that either keeping or breaking promises has any intrinsic deontic value. On his view, the values that promise keeping and promise breaking might have only arise from their being the exercise of a virtue or vice. If one holds that keeping or breaking promises have intrinsic deontic moral significance, then Conee's proposal will not be acceptable. See Conee, "The Moral Value of Promises," *Philosophical Review* 109 (2000): 411–22.

subsequent duty for himself. Similarly a driver might back her truck into a car, thus creating a duty to compensate the car's owner. Even if she would carry out this subsequent duty, we don't conclude from this that she has a *prima facie* duty to back her truck into the car and create the subsequent duty for herself. The lesson from these cases is that there is no positive value to satisfying a created duty, but only negative value to violating such a duty.

Just as making a promise, impregnating a woman, or backing one's truck into a car create new duties, so, within a subjective deontological code, acquiring morally relevant information creates a new duty. For example, let's suppose that if Claire investigates, she will come to believe of Mina that she is the least productive employee and so will come to have a *prima facie* duty to do what she believes to be laying Mina off. On the other hand, if she doesn't investigate, she will not believe of any employee that he or she is least productive, and so will come to have a *prima facie* duty to do what she believes to be selecting the most recently hired employee for layoff.²⁷ Thus, according to a subjective deontic moral code, gathering information (or failing to do so) can create subsequent duties for the agent because gathering information produces belief-states that generate new duties. These created duties must be handled in the same way that the duties created by making promises are: fulfilling such *prima facie* duties must be accorded no positive value, even though violating these duties is accorded negative value.

When we accept this, Claire's choices should then be represented as having revised deontic values, as shown in figure 2.²⁸ Figure 2 incorporates the new assumption that no fulfillment of a duty receives any positive value. The stars in figure 2 indicate that if Claire investigates, she will rightly lay off whichever employee she discovers to be the least productive, for a value of 0, while if she doesn't investigate, she will then lay off Max, whom she believes to be most recently hired, also for a value of 0. Thus Claire's doing what she believes is gathering information and her doing what she believes to be not gathering information would have the same value (0), and, according to Code C, it is not morally obligatory for her

27. In this particular case of not investigating, it may be more accurate to say that not investigating maintains, rather than produces, the lack of a belief about any employee that he or she is least productive. However, since Claire has an alternative (investigating) that would have eliminated this epistemic state, not investigating should be understood as playing an active role.

28. I represent the fact that an act of fulfilling a created duty receives "no positive value" by assigning it a value of 0. This number, along with other value numbers ascribed to acts, could be interpreted as measures on an interval scale, in which case the 0 is arbitrary. The point is not that we ascribe any particular numerical value, such as 0, to the act that fulfills a created duty, but rather that our ascribed value to this act is the same value it would receive if it did not constitute the fulfillment of this duty.




FIG. 2

to gather information, since she believes that she would achieve the same deontic value whether she investigates or not. This flies in the face of our assumptions about what it is morally incumbent on Claire to do, even according to a subjective code. Cases like Devon's compel us to adopt a new way of assigning deontic values to the fulfillment of duties, but this results in the completely counterintuitive conclusion that Claire has no subjective duty to investigate. This is hardly the result that subjectivizing deontologists hoped for.

On this new way of calculating the relevant deontic values, we do obtain a different prescription if we assume that Claire would fail to carry out her subsequent obligations. Thus suppose Claire dislikes Mina. If she investigates she will lay off whichever employee she finds to be least productive, but if she doesn't investigate she will bend the rules by laying off Mina rather than doing what she believes to be laying off the most recently hired employee. Thus investigating would have a deontic value of 0, but not investigating, which would lead her to violate her duty, would have a deontic value of -1. Assuming she foresees all of this, she has an obligation to investigate. However, this kind of effect only occurs when the agent would fail to carry out her subsequent duty. In cases having the structure of Claire's case, if the agent would carry out her subsequent duties, she has no duty to investigate. It is hardly a sat-

isfactory account of the duty to investigate if it obtains only when one would not carry out one's subsequent duty if one failed to investigate.

To complete this picture it is important to point out that some cases will include an independent reason to investigate—independent of the deontic value the agent believes that investigating would subsequently produce. In these cases things may turn out differently than they do when the only source of a duty to investigate is the derivative duty. For example, Claire might have signed a contract requiring her, as the company's human resource manager, to investigate employee productivity records before conducting layoffs. In this case Claire would have a "backwards-looking" duty to gather information, based on her contractual promise.²⁹ But in the original case there is no such independent duty to investigate: in that case, whether or not Claire has an information-gathering duty depends solely on whether doing so would produce the maximum amount of subjectively expected deontic value.

It may be tempting to suspect that Claire's case is unusual. If so, any general argument that subjective deontological codes cannot show that agents have suitable duties to investigate would require us to canvass a great many such cases. However, without such canvassing, one can demonstrate that Claire's case is not unusual: if fulfilling a duty has no positive deontic value, then in any case in which the only moral reason to do what the agent believes is investigating is the derivative reason I have described, there will never be a duty to do what the agent believes to be investigating unless the agent believes that investigating would lead to her failing to fulfill a subsequent duty.

Here is an informal proof: suppose the agent has two options, doing what she believes to be investigating (*A*) and doing what she believes to be not investigating (*B*).³⁰ She believes that if she does *A*, she

29. Suppose her contract requires Claire to investigate employee productivity records before releasing any employee in a financial exigency. Remembering the contract, she believes she has such a duty. Since she created this duty for herself by signing the contract and so believing it is in force, the deontic value of fulfilling her duty to do what she believes is fulfilling her contract is 0, while we might stipulate that the deontic value of doing what she believes is violating it is -3 . Thus the deontic value of each sequence of acts following her doing what she believes to be not investigating (and so doing what she believes to be breaking her contract), would gain an additional disvalue of -3 , while her act of doing what she believes to be investigating (and so doing what she believes to be honoring her contract) would gain an "additional" deontic value of 0. Thus investigating would turn out to maximize deontic value (it would produce a total deontic value of 0, whereas her doing what she believes to be not investigating would produce a total deontic value of -4). Hence, because of this independent reason, Claire would have a duty to do what she believes to be investigating.

30. Of course Claire's options might be more complex. For example, she might have different ways of investigating: she could read the company's personnel records, she could consult her assistant manager, or she could observe Max and Mina for an hour to see which

will then believe she has three (or more) options, *X*, *Y*, and *Z*, and also that she will come to believe of one act (say *X*) that its features are deontically better than those of the others, so *X* will be obligatory. According to our assumption that fulfilling a duty has no positive value, the deontic value of *X* is 0.

Assume next that the agent believes that if she does *B* (not investigating), she will then believe she will have three (or more) options, *K*, *L*, and *M*, and also that she will come to believe of one act (say *K*) that its features are deontically better than those of the others, so *K* will be obligatory. According to our assumption that fulfilling a duty has no positive value, the deontic value of *K* is 0.

Assume that the agent believes that if she does *A* she will carry out *X*, which she then has a duty to do. Assume also that the agent believes that if she does *B* she will carry out *K*, which she then has a duty to do. Since both *X* and *K* have a deontic value of 0, the agent has no derivative duty to do *A* (investigate), since it does not lead to greater deontic value. (This proof can be extended to cases in which several of the best actions subsequent to *A*, e.g., *X* and *Y*, have equally high deontic value, if we assume, plausibly, that merely permissible actions have no positive deontic value.) Thus in any case having a similar structure to Claire's case, and in which the agent will fulfill any duties she believes herself to have, the agent will have no duty to seek additional information before acting.

IV. TYING UP LOOSE ENDS

To complete the argument, let's return to tie up some loose ends.

A. Extension of Principle SD to Cases of Uncertainty

I noted earlier that, for simplicity, I would focus primarily on cases involving agents who have false beliefs or no beliefs rather than uncertain beliefs about their prospective actions. Claire's case fits this description; while some of her beliefs are true, others are false or simply lacking altogether. However, most agents have beliefs with varying degrees of certainty and uncertainty. Thus Claire, instead of being certain at t_1 that Max is the most recently hired employee, might think there is a .9 chance that Max was hired most recently, or might think there is only a .8 chance that she will fire the person she identifies as the least-

one seems least productive. These acts might have their own prima facie values (observing Max and Mina might humiliate them), and these would have to be taken into account in deciding what Claire's best overall course of action is. The proof assumes that we are focused only on the derivative values arising from Claire's epistemic state after she investigates or fails to investigate.

productive employee. Principle SD, together with the prohibition against assigning positive deontic value to fulfilling an obligation, must be extended to cover such cases as well. But doing so involves resolving dauntingly complex theoretical issues. For example, how would one restate Code C to apply to cases in which the agent is uncertain about the relevant facts? This seems clear enough for clause 2: it could be restated as “An agent subjectively ought *prima facie* to gather information before fulfilling her subsequent duties if and only if she believes that doing so would maximize subjective expected deontic value.” But what about clause 1? It makes no sense to import the concept of “subjective expected value” into clause 1, because the fulfillment (or violation) of clause 1 is what itself bears value. Clause 1 doesn’t ascribe a duty to the agent in virtue of something else her action might produce. Still, Claire could be uncertain whether laying off Mina would be laying off the least-productive employee. This suggests that we need a new code incorporating multiple clauses in place of the earlier clause 1. Each clause would define a distinct duty, such as a duty to “lay off a person who you think has a .9 chance of being least productive,” a duty to “lay off a person who you think has a .8 chance of being least productive,” and so forth. What is the deontic value of fulfilling these various duties? Perhaps, if laying off someone you think has a 1.0 chance of being least productive has a deontic value of 2, then laying off someone you think has a .9 chance of being least productive has a lesser deontic value of 1.8 ($2 \times .9$). But what is the deontic value of laying off someone you think has merely a .4 chance of being least productive? Should this be a small positive value, or a negative value? Moreover, if laying off someone you think has a 1.0 chance of being least productive has a value of 0 rather than 2 (as I have argued), then what is the deontic value of laying off someone you think has a .9 chance of least productive? Should this value also be 0, or should it be less than 0, or more than 0?

Another issue is how to handle an agent’s not having any belief about some matter. If Claire doesn’t investigate, should she be understood as thinking there is a .5 chance that Mina is least productive and a .5 chance that Max is least productive? This may be appropriate when Claire believes there is a definite number of prospective layoff candidates, so that she can partition her credences equally among them. But what if she doesn’t have any idea how many individuals are candidates for layoff? Some decision theorists have argued persuasively that in certain cases of this type there is no way for Claire to partition her credences so that she assigns equal probabilities to each one.³¹ All this makes exploration of how to extend principle SD to probabilistic cases highly com-

31. See Alan Hajek, “What Conditional Probability Could Not Be,” *Synthese* 137 (2003): 273–323, secs. 2.3 and 7.1.1.

plex because of the theoretical issues that need to be addressed. In light of these complexities, I shall defer extending principle SD to cases involving uncertain beliefs. Nonetheless it is clear that SD, together with the assumption that fulfilling an obligation has a deontic value of 0, generates an unacceptable answer to the question of whether Claire ought to investigate when she assigns a probability of 1.0 to the various possible upshots of her investigating and not investigating. Principle SD goes wrong in at least this central case, and its doing so raises doubts that it will go better in cases involving uncertainty.

B. Reasonable Belief Versions of the Subjective Duty to Investigate

What would be said about Claire's obligation by a theorist who believes "subjective" obligations are determined by what it would be reasonable for the agent to believe, not by what the agent actually believes? Would switching to the "reasonable belief" account rescue the subjective deontic duty to investigate? To answer this, let's restate Code C as follows in a "reasonable belief" version:

Code C*

- (1) When an agent reasonably believes she has been tasked to lay off employees in light of a financial retrenchment, the agent *prima facie* ought subjectively either to
 - (a) do what she reasonably believes to be laying off an employee if she reasonably believes that person to be the least-productive employee, or
 - (b) do what she reasonably believes to be laying off the most recently hired employee if she does not reasonably believe of any employee that he or she is the least-productive employee.
- (2) When an agent reasonably believes she has the opportunity to gather information relevant to a subsequent duty, then the agent subjectively ought *prima facie* to do what she reasonably believes to be gathering relevant information before fulfilling her subsequent duties if and only if she reasonably believes that doing so would subsequently lead her to produce the maximum amount of deontic value.

The phrase "what it would be reasonable for the agent to believe" has been interpreted in a number of different ways. According to the most plausible (in this context) interpretation of this phrase, it refers to *what the agent would be justified in believing given the evidence actually accessible to her*. On this interpretation, Claire's beliefs at t_3 count as reasonable if she

investigates (and comes to believe that Mina is least productive). So do her beliefs at t_3 if she fails to investigate (and so, let us assume, reasonably believes that Max is most recently hired but doesn't believe of any employee that he or she is least productive). Again we can assume that Claire reasonably believes that she will carry out her duty as she perceives it at t_3 . Then, on this interpretation of "reasonable beliefs," Code C*, together with the view that fulfilling created duties itself has no positive value, would counterintuitively imply—like the actual belief view—that Claire has no obligation to do what she believes to be investigating, since investigating would not maximize deontic value. Thus invoking the reasonable belief interpretation of subjective deontological codes does not rescue them from the problem we have discovered.

C. Implications for Objective Deontological Codes

We should ask whether denying positive value to the fulfillment of duties leaves intact our account of the duty to gather information in the context of objective deontological moral theories. The answer is "yes." We said that an agent has a derivative objective duty to gather information if and only if doing so would maximize deontic value. The first thing to note is that in the case of an objective deontological code, the agent's subsequent duty (for example, Claire's duty to lay off the least productive employee) is not a created duty. Her investigating which employee is least productive does not create her duty to lay off Mina; she has this duty on Tuesday whether or not she has investigated.

I claimed only that it is created duties whose fulfillment should be accorded no positive value. Thus the bar on ascribing positive value doesn't apply in the case of fulfilling objective duties. However, it is worth exploring what the upshot would be if we did import this bar to the case of non-created duties. If we take this course we can understand Claire's case as follows, interpreted as a case in which she has an objective duty of justice to lay off the least productive employee (regardless of what she happens to believe about the case). Code C, restated as an objective code, becomes the following Code C#:

Code C#

- (1) When an agent has been tasked to lay off employees in light of a financial retrenchment, the agent *prima facie* ought objectively either to
 - (a) lay off an employee if that person is the least-productive employee, or
 - (b) lay off the most recently hired employee if there is no unique least-productive employee.

- (2) When an agent has the opportunity to gather information relevant to a subsequent duty, then the agent objectively ought *prima facie* to gather relevant information before fulfilling her subsequent duties if and only if doing so would subsequently lead her to produce the maximum amount of deontic value.

Note that the “lesser” duty to lay off the most recently hired employee only kicks in—that is, an agent only counts as fulfilling or violating this duty—in cases in which there is no least productive employee. Assume that Mina is the least-productive employee but Max is the most recently hired employee. Claire’s objective options are represented in figure 3.

Assume that if Claire investigates, she will discover that Mina has been least productive and so will lay her off, for a deontic value of 0. Assume also that if she fails to investigate, she will continue to be uncertain which employee is least productive, but luckily will lay off Mina, the least productive employee (0). Since Claire would fulfill her layoff duty either way, she has no duty to investigate. This is consonant with an objective code, since she fulfills the same duty at t_2 in any case—unlike the situation with the subjective code, when her duty (and its importance) varies depending whether she investigates or not.

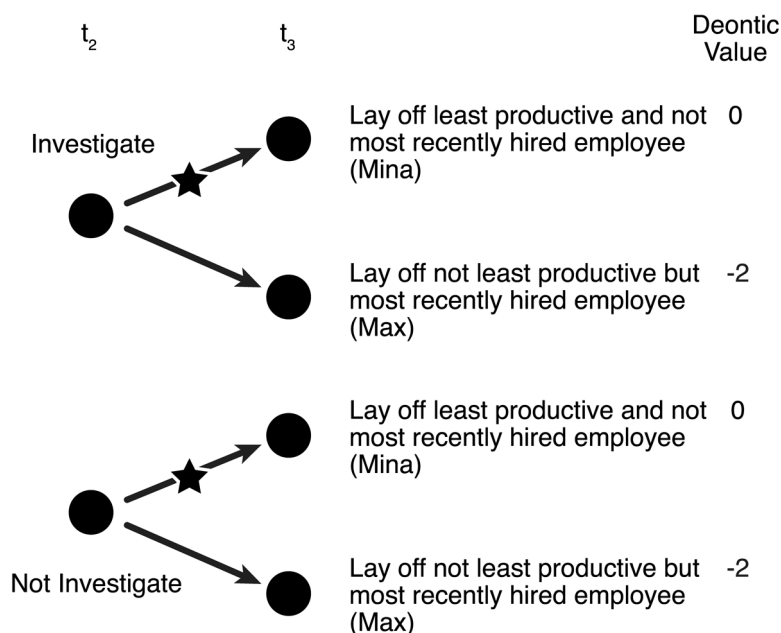


FIG. 3

Why are objective deontic codes immune to the problem we have discovered for subjective codes? The difference arises precisely because, according to the objective code, Claire ought to lay off the least productive employee, regardless of whether she has investigated or not and regardless of what she believes. In contrast to this, her subjective duty depends on what she believes, which in turn depends on whether or not she has investigated. On the subjective code, if she hasn't investigated and doesn't believe of any employee that he or she is least productive, then her duty is different. Duties under the objective code are not dependent on beliefs in this fashion.

We can conclude that even when we import the unargued-for assumption that fulfilling a noncreated duty receives no positive value, our recommended approach to determining a *prima facie* objective deontic duty to investigate still provides appropriate answers. If the agent would luckily fulfill the same subsequent duty whether she investigates or not, then both investigating and not investigating are permissible. If she would fulfill her subsequent duty only if she investigates, she has a duty to investigate. If investigating would not lead to her fulfilling her subsequent duty, she will not have a duty to investigate. These are appropriate outcomes for an objective code.

V. CONCLUSION

Because decision makers frequently lack the information necessary to use standard objective moral theories for guidance in making moral decisions, many theorists have proposed rejecting these theories in favor of subjective theories that ascribe moral status to actions as a function of what the agent believes about her prospective actions' circumstances and consequences, rather than as a function of the actions' actual circumstances and consequences. Some theorists, including Prichard, who famously advocates this "subjective turn" in normative theory, have asked whether such subjective theories can properly underwrite a duty to gather information before making a morally momentous choice.

We briefly noted that subjective welfare-maximizing codes can incorporate a normatively appropriate duty regarding information gathering.³² But the story is different with subjective codes that include "deontic" duties to do what one believes to be keeping promises, compensating others for past wrongs, and acting justly. We considered two

32. Ascribing no positive value to fulfilling duties does not affect subjective consequentialist theories, since on these theories the rightness or wrongness of investigating depends, not on the deontic value of the agent's subsequent fulfillment or violation of duties, but rather on the value of the events and states of affairs (such as effects on human welfare) that she believes her investigating would lead to.

ways to incorporate an information-gathering duty into subjective deontological moral codes: (1) as a freestanding duty characterized without reference to what the agent believes about how carrying out that duty would affect whether she carried out her subsequent deontic duties and (2) as a duty to do what the agent believes is gathering information if and only if she believes that doing so would maximize the deontic value of her subsequent actions. I argued that the second of these two approaches is the correct one, but that it fails to deliver what subjective deontological theories need: a normatively plausible duty to do what an agent believes is gathering information.

If subjective deontological moral codes are to provide a satisfactory account of a duty to gather information, the task remains to be accomplished. Absent such an account, it appears that such subjective codes fail to define adequately, or actually impose on us, an important type of duty that we all believe we have. Theories failing to impose such duties are crucially defective. Prichard may have been right to fear that this could be a fatal flaw in a subjective deontological theory.