Synthese (2014) 191:2835–2847
DOI 10.1007/s11229-014-0423-4

Deontology and doxastic control

Nicholas Tebben

Received: 13 September 2013 / Accepted: 14 February 2014 / Published online: 4 March 2014

Abstract Matthias Steup has developed a compatibilist account of doxastic control, according to which one's beliefs are under one's control if and only if they have a "good" causal history. Paradigmatically good causal histories include being caused to believe what one's evidence indicates, whereas bad ones include those that indicate that the believer is blatantly irrational or mentally ill. I argue that if this is the only kind of control that we have over our beliefs, then our beliefs are not properly subject to epistemic evaluation in deontological terms. I take as premises the claims (1) that acts which violate a deontic standard must be under the control of the agent that performs them, and (2) that deontic standards *are* deontic standards only if there is both something that it is to comply with them, and something that it is to violate them. The argument proceeds by showing that any belief which one might take to violate a deontic standard of a distinctively epistemic kind has a "bad" causal history, and so is, according to the compatibilist account, not under our control. Since these beliefs are not under our control, it follows from premise (1) that they do not violate any deontic standards of a distinctively epistemic kind. It then follows, from premise (2), that there are no deontic standards, of a distinctively epistemic kind, that govern belief. So if we have only compatibilist control over our beliefs, our beliefs are not properly subject to epistemic evaluation in deontological terms.

Keywords Doxastic voluntarism · Epistemic deontology · Epistemic justification · Freedom of belief

N. Tebben (✉)
Department of Philosophy & Religious Studies, Towson University, 8000 York Rd., Towson, MD, USA
e-mail: ntebben@towson.edu

1 Introduction

In a famous paper from the late 1980s, William Alston argued that we enjoy only very limited control over our beliefs (see Alston 1988). He surveyed three ways in which we might control our beliefs: directly, indirectly, and in a long range way. All three are found wanting. We have, he says, direct control over whether or not we raise our arms: simply deciding to raise one's arms is (usually) sufficient to raise them. But we have no such control over our beliefs. Alston challenges his readers to believe, simply by deciding to, that the United States is still a colony of England, and predicts that they will be unable to do so. He also argues that we lack indirect control over our beliefs. Indirect control is the kind of control that we enjoy over whether or not the light is on: turning the light on or off requires more than a decision, it also requires flipping a switch. But, just as you cannot believe that the US is still a colony simply by deciding to do so, you cannot flip a switch and believe it either.

Finally, there is long range control. Exercising long range control requires carrying out a plan over a long stretch of time, which is at least sometimes interrupted by doing other things. Whether or not one writes a book is something over which one has long range control. Although Alston admits that it is possible to exercise long range control over what one believes, he denies that our attempts to exercise long range control over our beliefs reliably succeed. "*Sometimes* people succeed" he says "in getting themselves to believe (disbelieve) something. But I doubt that the success rate is substantial. To my knowledge there are no statistics on this, but I would be very much surprised if attempts of this [long range] sort bore fruit in more than a small proportion of cases." (Alston 1988, p. 276) He then goes on to provide reasons to support his conjecture.

Alston's argument that we enjoy only very limited control over our beliefs was the first step in a larger argument. Epistemic justification is naturally thought of in terms of *blamelessness*: those whose beliefs are justified are not blameworthy for holding them (Alston 1988, pp. 257–258). One natural way to understand blamelessness is in terms of comportment with one's obligations. So, this line of thinking goes, those who are epistemically justified in holding their beliefs are those whose epistemic behavior comports with their epistemic obligations. Alston is skeptical of the idea that there are epistemic obligations, or, more generally, that beliefs are subject to deontic appraisal. So his larger purpose was to argue that "despite the connotations of the term, we are ill advised to think of being epistemically justified in terms of freedom from blame for believing."[1] Although Alston's version of the argument is somewhat more complicated, the off-brand version of his argument is very simple: deontological concepts correctly apply only to those things over which one enjoys a substantial degree of control, we do not enjoy a substantial degree of control over our beliefs, hence deontological concepts do not correctly apply to our beliefs. Since they do not, it is better to think of epistemic justification in non-deontic terms (or, perhaps, to evaluate beliefs in terms of something other than justification entirely).

[1] Alston (1988), p. 294. I have omitted one of Alston's notes. Alston also suggests that there is a range of "epistemic desiderata", and that we do better to evaluate beliefs relative to the desiderata found within that range, rather than in terms of justification. See Alston (2005).

There is reason to suspect, however, that there are more kinds of control than those that Alston considers. Matthias Steup has recently tried to rehabilitate the idea that we have a substantial degree of control over our beliefs (see Steup 2012, p. 153). He distinguishes between libertarian control and compatibilist control. In the general case, libertarians hold that one has control over whether or not one ϕs if and only if, in one and the same set of circumstances, one could intentionally ϕ, and one could intentionally refrain from ϕing. Libertarians thus take control and causal determination to be incompatible, for if one is caused to ϕ, then it is not the case that one could (in the relevant sense of 'could') refrain from ϕing, and if one is caused to refrain from ϕing, then it is not the case that one could (again, in the relevant sense) ϕ. A libertarian conception of *doxastic* control is one according to which one has control over one's belief about whether p, if and only if, in one and the same set of circumstances, one could intentionally believe that p and one could intentionally refrain from believing that p.

Compatibilists, by contrast, deny that having control over whether or not one ϕs requires that one could ϕ, and that one could refrain from ϕing, in one and the same set of circumstances. On their view, having control over whether or not one ϕs is, therefore, compatible with having been causally determined to ϕ, and with having been causally determined to refrain from ϕing, as the case may be. As our volitions play an important role in the production of our actions, compatibilists also maintain that our volitions may be under our control, even if they are causally determined. As Steup puts the point:

> Here is the key thought to which compatibilists are committed: Causal determination is not an obstacle to being in control of one's will. What matters is not *that* one's volitions are caused but *how* they are caused. Volitions with a good causal history are those that are under one's control; volitions originating in bad causes are those that fail to be under one's control. Thus, when compatibilists look at a particular action and the volition from which it flows and wish to judge whether the action was under the agent's control, they will have to determine whether the causes leading to the agent's volition and subsequent action were good or bad. (Steup 2012, p. 153.)

Steup takes it that beliefs can be formed on the basis of volitions, which are shaped by one's epistemic reasons. (See Steup 2012, p. 157 and the discussion that follows.)[2] A compatibilist conception of *doxastic* control, then, is one according to which beliefs with "good" causal histories are those over which one has control, and those with "bad" causal histories are those over which one does not have control. Steup argues that there is a strong prima facie case that we enjoy the kind of control over our beliefs that is captured by the compatibilist conception of doxastic control.

My purpose here is not to argue against Steup. Indeed, I suspect that he is right, and that we do enjoy compatibilist control over our beliefs. But I would like articulate a challenge that the view must overcome. And it is a challenge that suggests that

[2] Buckareff, however, argues that there are important differences between doxastic decisions and practical decisions, which pose severe problems for Steup's account of doxastic control. See Buckareff (2006). But see also Steup (2008, 2011a), for ways in which Steup would deal with these difficulties.

compatibilist control, at least in the doxastic case, is not an especially important kind of control. As Alston recognized, what *matters* about doxastic voluntarism is its close connection to epistemic deontology. Now, a central component of a compatibilist account of doxastic control is a specification of the conditions that distinguish good causal histories from bad ones. One of the main tasks for a doxastic compatibilist is, then, to draw this distinction in a way that permits the evaluation of beliefs in deontological terms. I will argue that this task is not trivial, indeed, I will argue that what is, perhaps, the most natural way of drawing the distinction—and one suggested by recent work on doxastic compatibilism—does not permit the deontological evaluation of beliefs.

The problem, I will argue, is that compatibilism, if it is developed in this natural way, does not allow for the violation of doxastic deontic standards. But any genuine deontic standard must be such that it can be violated. So, absent a different way to mark out the distinction between good and bad causal histories, if we enjoy only compatibilist control over our beliefs, beliefs are not governed by deontic standards. If that is the case, we would do better to think of epistemic justification in other terms (or to stop evaluating beliefs in terms of their justificatory status entirely). Alston's argument against a deontic construal of epistemic justification turned on denying doxastic voluntarism, but the same result can be had if the only kind of doxastic control that we enjoy is compatibilist control of this sort. In section three I will consider other ways to develop compatibilism, and provide some considerations that suggest that they are also problematic.

2 The argument

2.1 A point about normative standards

Whatever else they may do, normative standards place a partition on the class of things (or acts, or people, or whatever, both actual and merely possible) to which they apply, sorting them into a class of things (or acts, or etc.) that do well by that standard, and a class that does not. (They may also, of course, separate those classes into still finer classes. "Doing well" by a normative standard may be a matter of degree.) Think about assigning grades to students. In ordinary classes, assigning an A to a student is a means of *evaluating* him or her positively. There are normative standards of behavior in the class, requiring, for example, providing correct answers on the exams. The assignment of an A can function as an evaluation only because it reflects the result of measuring a student's performance against the standards of the class, and finding that the student did well instead of poorly.

Now imagine a teacher who necessarily gives each of the students in her class an A. In this teacher's class, A's no longer function as a means of evaluation. An A in such a class might have the same function as a certificate of participation, providing an indication that the student was enrolled in the class, but it would not reflect how well or poorly the student did relative to the standards of the class. If, necessarily, everyone receives an A, in giving A's we will not be praising all students equally, we will have stopped praising students at all. Similar remarks apply to negative evaluations. An F

is a mark of failure only because there is the possibility of not failing. If a teacher necessarily gives everyone an F, receiving an F is not a way of being evaluated.

The point that this example illustrates is a general one about normative standards. There is something that it is to do well by a normative standard only if there is something that it is to do poorly, by that standard. And, likewise, there is something that it is to do poorly by a standard only if there is something that it is to do well by that standard. If the class of things (or people, or whatever) that do poorly by a normative standard is necessarily empty, or if the class that does well is necessarily empty, then so is the other. Doing well requires the possibility of doing poorly, and doing poorly requires the possibility of doing well.

Deontic standards are a species of normative standard, and inherit this feature from the genus. Those who comport with their duties do *well*, relative to their obligations, and this is possible only if there is something that it is to do poorly. Likewise, those who violate their duties do *poorly*, relative to their obligations, and this is possible only if there is something that it is to do well. We could have no obligations which are such that, necessarily, everyone violated them, nor could we have obligations which are such that, necessarily, everyone discharges them. We will say that this means that, for any deontic standard, there must be both something that it is to comply with it, and something that it is to violate it.

It is worth distinguishing this claim from the ought-implies-can principle (hereafter: OIC). Assume that one ought to discharge one's obligations. OIC, then, implies that it is possible to discharge one's obligations, which, in turn, implies that there is something that it is to discharge those obligations. My claim, on the other hand, implies that there is something that it is to discharge one's obligations, but not that it is possible to discharge them. Even if there is something that it is to discharge one's obligations, it may not be possible to do so, given, say, one's practical circumstances.

2.2 Compatibilism

Recall that compatibilists take the quality of the cause of a volition to determine whether or not it is under an agent's control. Elaborating on this point Steup says:

> If a volition's causal history is a good one – indicating a rational and healthy mind – then the volition is one that was under the agent's control. If, on the other hand, a volition's causal history is a bad one – indicating blatant irrationality or even mental illness – then the volition was not under the agent's control. In either case, the agent cannot decide otherwise under the same circumstances, and it is not the case that two distinct futures are open to the agent. (Steup 2012, p. 147.)

We said above that an essential element of a compatibilist theory is a specification of the conditions that distinguish good causal histories from bad ones. This passage does not provide a full characterization of the relevant distinction—some causal histories indicate neither rationality nor irrationality, and neither health nor illness—but it is an intuitive and appealing start to a compatibilist theory. We do seem to be in control when our minds are functioning well, and health and rationality seem to be indications of a well-functioning mind. Likewise, we do not seem to be in control when our

minds are functioning poorly, and blatant irrationality and mental illness seem to be indications of a mind that is functioning poorly. I will argue, however, that any theory that distinguishes good from bad causal histories in the way suggested here will not provide the grounds for the application of deontic concepts to beliefs.

The problem is that those beliefs which violate our epistemic obligations (or otherwise violate deontic constraints of a distinctively epistemic nature) are a sub-set of those with "bad" causal histories. Indeed, when we violate our epistemic obligations, it is precisely because our beliefs have bad causal histories. Consider, for example, beliefs formed on the basis of the gambler's fallacy. Those who commit the gambler's fallacy take the fact that a system has produced one particular outcome in a relatively long string of cases to be grounds for inferring that the system is, therefore, less likely to continue producing that outcome than it would be otherwise.

Say that Cole is playing roulette, and he observes the ball land on black ten times in a row. Cole then reasons that red is "due", and therefore believes that the odds that the ball will land on red in the next spin is greater than 1:2. This is a paradigmatic example of an unjustified belief. If there are any deontic constraints on belief, we should expect them to be in operation here: one is prohibited from forming beliefs on the basis of the gambler's fallacy.

Notice, however, that the reason that Cole is in violation of a duty (if he is), is that his belief is blatantly irrational. The gambler's fallacy is simply counter-induction (reasoning that the future will not be like the past), and reasoning counter-inductively is obviously irrational. What Cole *should* do is determine whether there is a greater than 1:1024 chance that the game has been rigged; if there is, he should believe that the odds that the ball will land on red are *less* than 1:2, because the fix is in. (Otherwise, he should believe that the odds are 1:2.)

Given Cole's evidence, if beliefs are subject to any deontic constraint, he is prohibited from believing that the odds that the ball will land on black are greater than 1:2. He violates a duty if he does so believe. But the reason that he is prohibited from holding this belief, the reason that he violates a duty in holding it, is precisely because it is blatantly irrational.

Consider another example. Say that Morrey examines all of the members of the L.A. Lakers basketball team and concludes that a sizable percentage of the people in L.A. are at least seven feet tall. Clearly this belief violates an epistemic obligation, if any do. But, again, the fact that it is forbidden is a product of the fact that it is blatantly irrational. Drawing hasty generalizations from biased samples is a paradigmatically irrational way of reasoning, and the fact that it is explains why one is forbidden from holding beliefs that are formed in this way.

Now, one important fact about deontic standards is that they are applicable only to behavior that is under one's control. Consider the fact that sneezing, something over which we enjoy only very limited control, is not plausibly subject to deontic evaluation. And the reason that it is not is not merely that, in most situations, nothing much turns on whether or not one sneezes. Imagine a situation in which quite a lot turns on whether or not one sneezes (perhaps one's family is hiding from a murderer bent on doing them in); though it would be *better* if one did not sneeze, even so, one violates no duties by so doing. The reason that sneezing violates no duties is that sneezing is not (usually) under one's control. We needn't beg any questions here; plausibly if we

enjoyed even compatibilist control over whether or not we sneeze, sneezing might be contrary to our duty. But we do not enjoy even compatibilist control over whether or not we sneeze, we usually enjoy no control at all over it, and hence it falls outside the range of deontic evaluation. If the compatibilist is right, Cole and Morrey enjoy no control over their beliefs, and, hence, their beliefs also fall outside of the range of deontic evaluation.

It might be tolerable if we were to find that those who commit the gambler's fallacy, or who generalize from biased samples, have no control over their beliefs, and so violate no epistemic obligations. It runs counter to our intuitive judgments, but this is not a conclusive objection. That intuitive judgments must be revised in the light of mature theory is not a new idea. But perhaps the lesson from these cases generalizes. Perhaps *all* beliefs which violate epistemic[3] deontic constraints (if there are any) are blatantly irrational. If that were the case, then we would have no control over *any* beliefs that violate deontic constraints, and so all "violations" of a deontic constraint would be merely apparent violations.

Now, whatever, precisely, epistemic rationality comes to, epistemic deontological constraints are constraints against violations of epistemic rationality. We might split the class of beliefs that violate standards of epistemic rationality into those that are blatantly irrational and those that are subtly irrational. Blatantly irrational beliefs are not under our control, by compatibilist standards. What about subtly irrational beliefs?

Nothing in the compatibilist position, as so far developed, closes off the possibility that one has enough control over them for these beliefs to be under our control. If they are, *and if subtly irrational beliefs violate deontic constraints*, then doxastic compatibilism leaves room for both compliance with, and violation of, deontic standards.

I would like to argue, however, that subtly irrational beliefs violate no deontic standards of belief. Assume, for the purposes of a *reductio*, that there are constraints against subtly irrational beliefs. Say that the belief that p is subtly irrational given evidence e; perhaps believing that p on this evidence involves a fallacy that ordinary people are unable to see. Now, if ought implies can (an assumption that I will make here), then if one ought to refrain from believing that p, it follows that one can refrain from believing that p. What modality is at issue here? It must be a very weak one. The principle *ought implies can* says more than that if one ought to do x it is logically possible for one to do x. It must be something like, if one ought to do x, then it is practically feasible for one to do x.[4]

Now, it is widely admitted (at least by both Alston and Steup) that, ordinarily, appreciation of one's evidence automatically determines one's beliefs.[5] If people are unable to see the fallacy involved in reasoning from e to p, then they will automatically

[3] This qualification will usually be omitted below. Understand that all normative standards and constraints mentioned in this section and the next are of a distinctively epistemic (as opposed to, say, moral, or legal) kind.

[4] Mason (2003) argues that the modality involved in the OIC principle must be such that agents are able to *intentionally* discharge their obligations. That would suffice for my purposes. If an agent is unable to notice a flaw in his or her reasoning, that agent is not able to intentionally avoid reasoning badly in that way.

[5] See, for example, Alston (1988) and Steup (2012). Steup, of course, does not take this fact to show that we are not in control of our beliefs.

believe that p, given e. Given that, it is not practically feasible for them to refrain from believing that p, given evidence e. Therefore, if ought implies can, it follows that it is not the case that they ought[6] to refrain from believing that p, even if a fallacy is involved in reasoning from e to p. Now, if subtly irrational beliefs are those that are produced by fallacious inferences that ordinary people are unable to detect, then, when presented with the relevant evidence, epistemic agents are ordinarily unable (in the relevant sense) to avoid holding those beliefs. If they are unable to avoid holding these beliefs, and if OIC is true, then it follows that it is not the case that they ought not to hold these beliefs. We can, therefore, maintain, on the assumption that OIC is true, that subtly irrational beliefs violate no deontic standards.

So for all x, if x violates a deontic standard of belief, then it follows that x is not subtly, but *blatantly*, irrational. This is not to deny that people often violate deontic standards of belief. Empirical research even suggests that we are naturally inclined towards committing the gambler's fallacy. (See Tversky and Kahneman 1974, p. 1125.) Even so, it is not hard to see that the gambler's fallacy is a fallacy, despite our frequent failures to see that it is. It is an instance of counter-inductive reasoning, and counter-induction is obviously a poor reasoning strategy.[7]

If every belief that violates a deontic constraint is obviously irrational, then they all have "bad" causal histories, involving, for example, counter-induction or generalization from a biased sample. But if they all have bad causal histories, then they are, by compatibilist lights, not under our control. Since all violations of a deontic constraint must be under one's control, it follows that, if compatibilism is correct, there is nothing that it is to violate an epistemological deontic constraint. But there *are* deontic constraints only if there is both something that it is to comply with them, and if there is something that it is to violate them. So, if we enjoy no control besides compatibilist control, it follows that there *are* no deontic standards applicable to beliefs. That is, if the compatibilist is right, and if we have no other kind of control, then the conclusion of Alston's argument is true, even if Alston was wrong about whether or not we enjoy a substantial degree of control over our beliefs.

3 Reason-responsiveness

The suggestion that blatantly irrational beliefs have bad causal histories is a natural and appealing one, since blatant irrationality seems to indicate a poorly functioning mind, and minds that are functioning poorly seem to be good candidates for those that are not under our control. We have seen, however, that this approach will not provide us with the resources to make sense of epistemic deontology.

[6] There is, of course, *some* sense in which one ought to refrain from going in for fallacious inferences, at least in that non-fallacious reasoning is (*qua* reasoning) better than fallacious reasoning. But this is an "external" sense of 'ought', and not one that has much bearing on the questions of epistemic deontology that are at issue here. That the reasoning is fallacious gives one a reason not to go in for it, but this is not a reason *for* the agent, not one that the agent possesses.

[7] Moreover, everyone recognizes that it is a poor reasoning strategy. We might reason counter-inductively when playing roulette, but no one concludes that water will not quench their thirst in the future from the fact that it has quenched their thirst in the past.

There are, however, other tacks that the compatibilist can take. There is, for example, a rich philosophical tradition which sees freedom as intimately linked with responsiveness to reasons. This connection has been productively explored by Fischer and Ravizza (see their 1998), and Steup has suggested that it may be significant for our understanding of doxastic control. (See Steup 2008, p. 390.) The insight that there is a close connection may help us draw a sensible line between good and bad causal histories, which will allow us to make room for the violation of deontic constraints.

So here is a suggestion: perhaps beliefs have good causal histories if and only if they are a product of processes that are responsive to reasons.[8] Now, the notion of *responsiveness* calls for clarification. I will begin by considering two relatively simple ways in which this notion might be developed, and I will argue that neither of them is satisfactory. A more nuanced account of reason-responsiveness can be found in Fischer and Ravizza's subtle and sophisticated theory of moral responsibility. This account is certainly an improvement over the simple ones. I will argue, however, that there are crucial differences between the moral and the epistemic cases, which make Fischer and Ravizza's theory ill-suited to serve as a model for a theory of doxastic control.

Consider first:

- Apt Evaluation: One's belief that p is responsive to reasons if and only if it was the product of an apt evaluation of one's epistemic reasons.

Apt Evaluation is the kind of reason responsiveness we make use of when we complain that those who cannot be talked out of their position won't listen to—are unresponsive to—reason. It is clear, however, that if we are to ground the application of deontic concepts to beliefs through our responsiveness to reasons, we must not construe *responsiveness* in this way. If we were to give it this construal, Cole would, again, come out as not being in control of his belief that the odds are greater than 1:2 that the ball will land on red. Notice that Cole's beliefs are not *appropriately* shaped by his reasons, and that this is exactly the reason that they are in violation of a deontic constraint.

The second account of reason-responsiveness that I would like to consider is one suggested by Steup. Steup is aware that accounts of reason-responsiveness must avoid the consequence that all irrational beliefs turn out not to be responsive to reasons. (He cites Bayer in connection with this point. See Steup (2011a), p. 547 n13.) In order to deal with this difficulty, he suggests that we adopt a theory of reason-responsiveness that allows reason-responsive agents to act against their reasons in a limited range of cases. He offers the following as an illustration of the kind of approach that he has in mind: "An agent's decision or choice to ϕ is caused in a reason responsive way iff in nearby worlds the agent tends to decide or choose to ϕ if and only if the totality of the agent's reasons supports ϕ-ing." (Steup 2011a, p. 547. I have omitted one of Steup's notes.) Put in an epistemic key this becomes:

- Likely Apt Evaluation: One's belief that p is responsive to reasons if and only if, in nearby possible worlds, one tends to believe that p if and only if the totality of one's epistemic reasons support believing that p.

[8] See Steup (2011b) for more information on this approach.

Steup offers this account as an illustration only; he does not claim that it is free from counter-examples. And indeed it is not. Say that Bart has been in a terrible auto accident, which has left him unable to draw inferences from his perceptual beliefs. In an effort to help him, the local hospital has created a device that will allow its operator to directly implant beliefs into Bart's mind while leaving the rest of his mind unaltered. With the intention of letting Bart live a normal life, they decide to implant in him all and only those beliefs that they take to be supported by his perceptual evidence. Let x be some belief that is supported by Bart's perceptual evidence. Now, Bart cannot infer that $x$is true, but since Bart's doctors take it to be supported by his evidence, the doctors push the button that causes Bart to believe that x.

Bart's belief that x meets the conditions on reason responsiveness captured in Likely Apt Evaluation. Bart's doctors would not have produced this belief in him if his evidence did not indicate that x, and whenever his evidence indicates that x, they would produce it in him. Nevertheless, there is reason to believe that it is not under his control. Bart's belief that x reliably varies with his reasons for belief, but Bart does nothing to help establish this connection. Bart's perceptions, and hence his reasons for belief, play no role in the production of his belief. It is, rather, the doctors' reasons for belief (in particular, their reasons for believing that Bart's reasons support x) that are causally responsible for Bart's belief. And so, it seems, the belief that x is not under Bart's control, it is under the control of his doctors.

If we are to sort good causal histories from bad ones on the basis of whether or not the belief-forming processes involved are responsive to reasons, we need an account of reason-responsiveness that, unlike Apt Evaluation, does not imply that irrational beliefs are not under our control, but which, unlike Likely Apt Evaluation, is not subject to counterexamples. Likely Apt Evaluation is modeled on Fischer and Ravizza's work, but its connection to their view is very loose. A more hopeful strategy would be to develop a theory that is modeled very closely on their account. However, caution is called for here. Their account of reason responsiveness is a part of an account of *moral* responsibility; I will argue that there are crucial differences between the moral and epistemic cases, such that Fischer and Ravizza's theory cannot simply be appropriated and employed for epistemological purposes.

Here is their summary of their account of reason responsiveness:

> A mechanism of kind K is ... responsive to reason to the extent that, holding fixed the operation of a K-type mechanism, the agent would *recognize* reasons ... in such a way as to give rise to an understandable pattern (from the viewpoint of a third party who understands the agent's values and beliefs), and would *react* to at least one sufficient reason to do otherwise (in some possible scenario).[Fischer and Ravizza (1998), pp. 243–244.][9]

One point that Fischer and Ravizza emphasize is that for one to be morally responsible for an action, it is not sufficient that that action merely have been the product of a reason-responsive mechanism. The mechanism must also be "one's own" (Fischer and

[9] This is a summary of their account of *moderate* reason responsiveness, which is the view that they ultimately settle on. They do identify other kinds of reason responsiveness, but they argue that these other kinds of responsiveness do not provide a satisfactory basis for moral responsibility.

Ravizza 1998, p. 207). Call the facts that must obtain in order for a reason-responsive mechanism to be ones own, its *ownership condition*. Some ownership condition is necessary for an account of moral responsibility because the features that constitute reason responsiveness could have been acquired in a way that undermines moral responsibility, for example, by direct stimulation of one's brain. (See Fischer and Ravizza 1998, p. 187.) It should be clear that any theory of doxastic control that is modeled on Fischer and Ravizza's view will also require an ownership condition, as direct stimulation of one's brain (as in the example involving Bart) should undermine control of belief as much as responsibility for action.

One virtue of a theory of doxastic control based on Fischer and Ravizza's work is that it leaves room for the violation of deontic constraints on belief. Given Cole's mistaken beliefs about probability, the reasons that he recognizes form a coherent pattern, but, nevertheless, there is *some* world in which he does not commit the gambler's fallacy (because in that world he has, say, taken a class on probability theory). So a theory based on Fischer and Ravizza's account of moral responsibility avoids the problem that we identified with Apt Evaluation. Whether or not it is free from counterexamples is a more difficult question, but, for present purposes, not an important one, because it is the *ownership* condition, not the account of reason-responsiveness itself, that poses a problem for epistemological adaptations of Fischer and Ravizza's view.

The ownership condition that Fischer and Ravizza develop has three components:

1. That one believe that one's "choices and actions are causally efficacious in the world" (Fischer and Ravizza 1998), p. 210),
2. That one believe "that [one] is a fair target of reactive attitudes as a result of how [one] exercises [one's] agency in certain contexts" (Fischer and Ravizza 1998, p. 211), and
3. That one holds these beliefs because of one's appreciation of the evidence that indicates that they are true. (Fischer and Ravizza 1998, p. 213.)

The ownership condition poses severe difficulties for those who would adapt Fischer and Ravizza's view for epistemological purposes. If this view is to be used for epistemological purposes, condition (1) must be taken to say that one must believe that one's *doxastic* choices and actions are causally efficacious. Now, an account of doxastic control is important to have because it would allow us to evaluate beliefs in deontological terms. I take it that a condition of adequacy on an account of doxastic control is that it permits such evaluation. That is not to deny that there are some whose beliefs are exempt from such evaluation: small children and those with serious mental illnesses, perhaps. But an adequate account of doxastic control must permit us to evaluate the beliefs of ordinary people in deontological terms.

But ordinary people have no beliefs about the efficacy of their doxastic choices; indeed, that there is such a thing as a doxastic choice is a controversial and abstruse philosophical thesis. So ordinary people do not meet the ownership condition, at least in the doxastic case. Since they do not meet the ownership condition, they do not have control over their beliefs, whether or not their belief-forming mechanism is reason-responsive. It follows, if doxastic control is to be modeled on Fischer and Ravizza's account of moral responsibility, that ordinary peoples' beliefs cannot be evaluated in

deontological terms. So such an account of doxastic control would fail to satisfy one of its adequacy conditions.

It must be admitted that there may be a different, and satisfactory, way to develop an ownership condition for reason-responsive mechanisms, but none are apparent to me. Other obvious candidates also require more self-reflection than most people go in for. One might, for example, require that one endorse the operation of one's own belief-forming mechanisms. But I suspect that relatively few people ever consider the possibility that they *have* belief-forming mechanisms. There are also problems involved with plausible counterfactual ownership conditions. For example, one plausible ownership condition is that one *would* believe that one's doxastic decisions are causally efficacious, if one were to consider the question. But it is far from clear what ordinary people would believe if they were to consider the question. After all, those philosophers who *do* consider the question are split on the issue. Suffice it to say that if there is an adequate account of doxastic control that is modeled on Fischer and Ravizza's theory of moral responsibility, an essential part of it would be an ownership condition that does not refer to believers' beliefs about their own doxastic choices or belief-forming mechanisms.

We began this section by considering two simple ways in which one might construe responsiveness to reasons. Neither was satisfactory. Apt Evaluation implies that all beliefs that violate deontic standards of belief are not under our control. Likely Apt Evaluation is subject to counterexample. We then considered the possibility that an account of doxastic control ought to be modeled closely on Fischer and Ravizza's account of moral responsibility. We saw, however, that there are important differences between the moral and the epistemic cases, such that Fischer and Ravizza's theory does not provide a satisfactory model for an account of doxastic control.[10] This survey of accounts of doxastic control in terms of reason responsiveness is not intended to be exhaustive, but it does indicate that there are severe challenges that any such account must face.

4 Conclusion

Steup suggests that we have control over those of our beliefs that have good causal histories, and that we do not have control over those with bad causal histories. An intuitive and attractive suggestion is that blatantly irrational beliefs are the product of bad causal chains, as blatant irrationality is a sign that one's mind is functioning improperly, and it is natural to think that one's beliefs are not under one's control when one's mind is not working well. And, indeed, that blatantly irrational beliefs have bad causal histories is suggested by recent work in doxastic compatibilism. But an account of doxastic control according to which all blatantly irrational beliefs have bad causal histories cannot ground the application of deontic concepts to beliefs. Those beliefs that are prohibited by deontic constraints on belief are precisely those that are blatantly irrational. So if this is the only kind of control over our beliefs that we enjoy, there are

[10] Of course nothing said here should be taken as criticism of Fischer and Ravizza, who did not intend to provide a model for an account of doxastic control.

no violations of epistemic deontic constraints that are under an agent's control. But behavior that is not under one's control cannot violate a deontic constraint. So if this is the only kind of control that we have, then there is nothing that it is to violate an epistemic deontic constraint.

Furthermore, deontic constraints are a species of normative standards, and something *is* a normative standard only if there is both something that it is to do well by it, and something that it is to do poorly by it. Since all violations – all believings that do poorly by the deontic standards of belief – are not under our control, according to the compatibilist's view, then if compatibilist control is the only kind of control that we enjoy over our beliefs, beliefs are not properly subject to deontic evaluation.

Another promising way to show that beliefs are under our control, in a compatibilist sense, is to take the distinction between good and bad causal histories to track the distinction between beliefs that are produced by reason-responsive processes and those that are not. Though this is a promising thought, obvious ways in which the notion of reason-responsiveness might be developed are unsatisfactory, and the sophisticated work of Fischer and Ravizza cannot be adapted for epistemological purposes in any straightforward way.

Acknowledgments I would like to thank John P. Waterman, Nick Goldberg, and two anonymous reviewers for their comments on earlier drafts of this paper. Earlier drafts were also presented to the 2013 meeting of the Alabama Philosophical Society and at the 2013 Northwest Philosophy Conference. I would like to thank my audiences at both conferences for their feedback.

References

Alston, W. (1988). The deontological conception of epistemic justification. In J. Tomberlin (Ed.), *Philosophical perspectives 2: Epistemology* (pp. 257–300). Atascadero, CA: Ridgeview.

Alston, W. (2005). *Beyond justification: Dimensions of epistemic evaluation*. Ithaca, NY: Cornell University Press.

Buckareff, A. (2006). Doxastic decisions and controlling belief. *Acta Analytica, 21*, 102–114.

Fischer, J. M., & Ravizza, M. (1998). *Responsibility and control*. Cambridge: Cambridge University Press.

Mason, E. (2003). Consequentialism and the 'ought implies can' principle. *American Philosophical Quarterly, 40*, 319–331.

Steup, M. (2008). Doxastic freedom. *Synthese, 161*, 375–392.

Steup, M. (2011a). Belief, voluntariness and intentionality. *Dialectica, 65*, 537–559.

Steup, M. (2011b). Justification, deontology and voluntary control. In S. Tolksdorf (Ed.), *Conceptions of knowledge* (pp. 461–486). Berlin: Walter de Gruyter.

Steup, M. (2012). Belief control and intentionality. *Synthese, 188*, 145–163.

Tolksdorf, S. (Ed.). (2011). *Conceptions of knowledge*. Berlin: Walter de Gruyter.

Tomberlin, E. (Ed.). (1988). *Philosophical perspectives 2: Epistemology*. Atascadero, CA: Ridgeview.

Tversky, A., & Kahneman, D. (1974). Judgment under uncertainty: Heuristics and biases. *Science, 185*, 1124–1131.